

Tensor Model Parallel Training

COMP4901Y

Binhang Yuan

Tensor Model Parallelism

Recall Data & Pipeline Parallelism

- Data Parallelism:
 - **Memory issue:** each device needs to maintain a complete copy of the model (parameters, gradients, and optimizer status).
 - **Statistical efficiency:** if the global batch size is too large, it may affect the convergence rate
- Pipeline Parallelism:
 - **Bubble overhead:** the pipeline parallelism efficiency decreases as the number of stages increases.

TransformerBlocks($x \in \mathbb{R}^{B \times L \times D}$) $\rightarrow x' \in \mathbb{R}^{B \times L \times D}$

- B is the batch size;
- L is the sequence length;
- D is the model dimension;
- Multi-head attention:
 $D = n_H \times H$
- H is the head dimension;
- n_h is the number of heads.

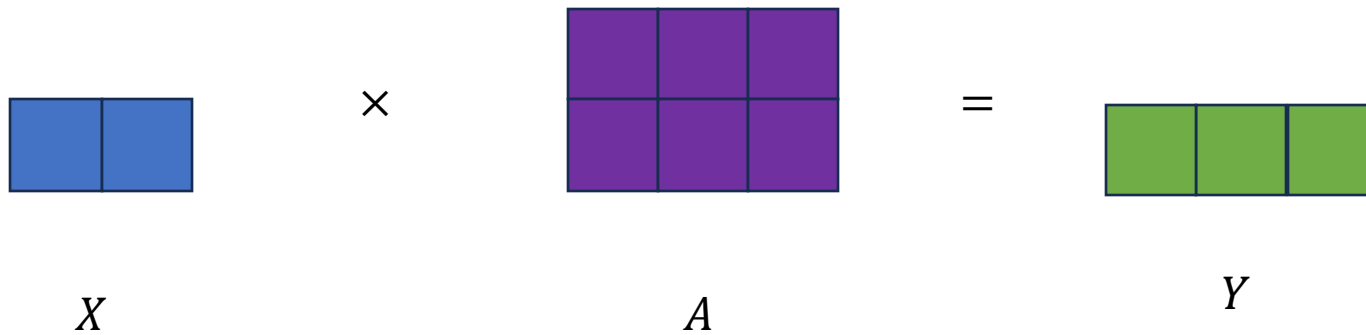
Computation	Input	Output
$Q = xW^Q$	$x \in \mathbb{R}^{B \times L \times D}, W^Q \in \mathbb{R}^{D \times D}$	$Q \in \mathbb{R}^{B \times L \times D}$
$K = xW^K$	$x \in \mathbb{R}^{B \times L \times D}, W^K \in \mathbb{R}^{D \times D}$	$K \in \mathbb{R}^{B \times L \times D}$
$V = xW^V$	$x \in \mathbb{R}^{B \times L \times D}, W^V \in \mathbb{R}^{D \times D}$	$V \in \mathbb{R}^{B \times L \times D}$
$[Q_1, Q_2 \dots, Q_{n_h}] = \text{Partition}_1(Q)$	$Q \in \mathbb{R}^{B \times L \times D}$	$Q_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, n_h$
$[K_1, K_2 \dots, K_{n_h}] = \text{Partition}_1(K)$	$K \in \mathbb{R}^{B \times L \times D}$	$K_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, n_h$
$[V_1, V_2 \dots, V_{n_h}] = \text{Partition}_1(V)$	$V \in \mathbb{R}^{B \times L \times D}$	$V_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, n_h$
$\text{Score}_i = \text{softmax}(\frac{Q_i K_i^T}{\sqrt{D}}), i = 1, \dots, n_h$	$Q_i, K_i \in \mathbb{R}^{B \times L \times H}$	$\text{score}_i \in \mathbb{R}^{B \times L \times L}$
$Z_i = \text{score}_i V_i, i = 1, \dots, n_h$	$\text{score}_i \in \mathbb{R}^{B \times L \times L}, V_i \in \mathbb{R}^{B \times L \times H}$	$Z_i \in \mathbb{R}^{B \times L \times H}$
$Z = \text{Merge}_1([Z_1, Z_2 \dots, Z_{n_h}])$	$Z_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, n_h$	$Z \in \mathbb{R}^{B \times L \times D}$
$\text{Out} = ZW^O$	$Z \in \mathbb{R}^{B \times L \times D}, W^O \in \mathbb{R}^{D \times D}$	$\text{Out} \in \mathbb{R}^{B \times L \times D}$
$A = \text{Out} W^1$	$\text{Out} \in \mathbb{R}^{B \times L \times D}, W^1 \in \mathbb{R}^{D \times 4D}$	$A \in \mathbb{R}^{B \times L \times 4D}$
$A' = \text{relu}(A)$	$A \in \mathbb{R}^{B \times L \times 4D}$	$A' \in \mathbb{R}^{B \times L \times 4D}$
$x' = A' W^2$	$A' \in \mathbb{R}^{B \times L \times 4D}, W^2 \in \mathbb{R}^{4D \times D}$	$x' \in \mathbb{R}^{B \times L \times D}$

Tensor Model Parallelism

- High-level idea:
 - The tensor is split up into multiple chunks;
 - Instead of having the whole tensor reside on a single GPU, each shard of the tensor resides on its designated GPU.
 - Each shard is processed separately and in parallel on different GPUs.
 - The results are synchronized at the end of the step.

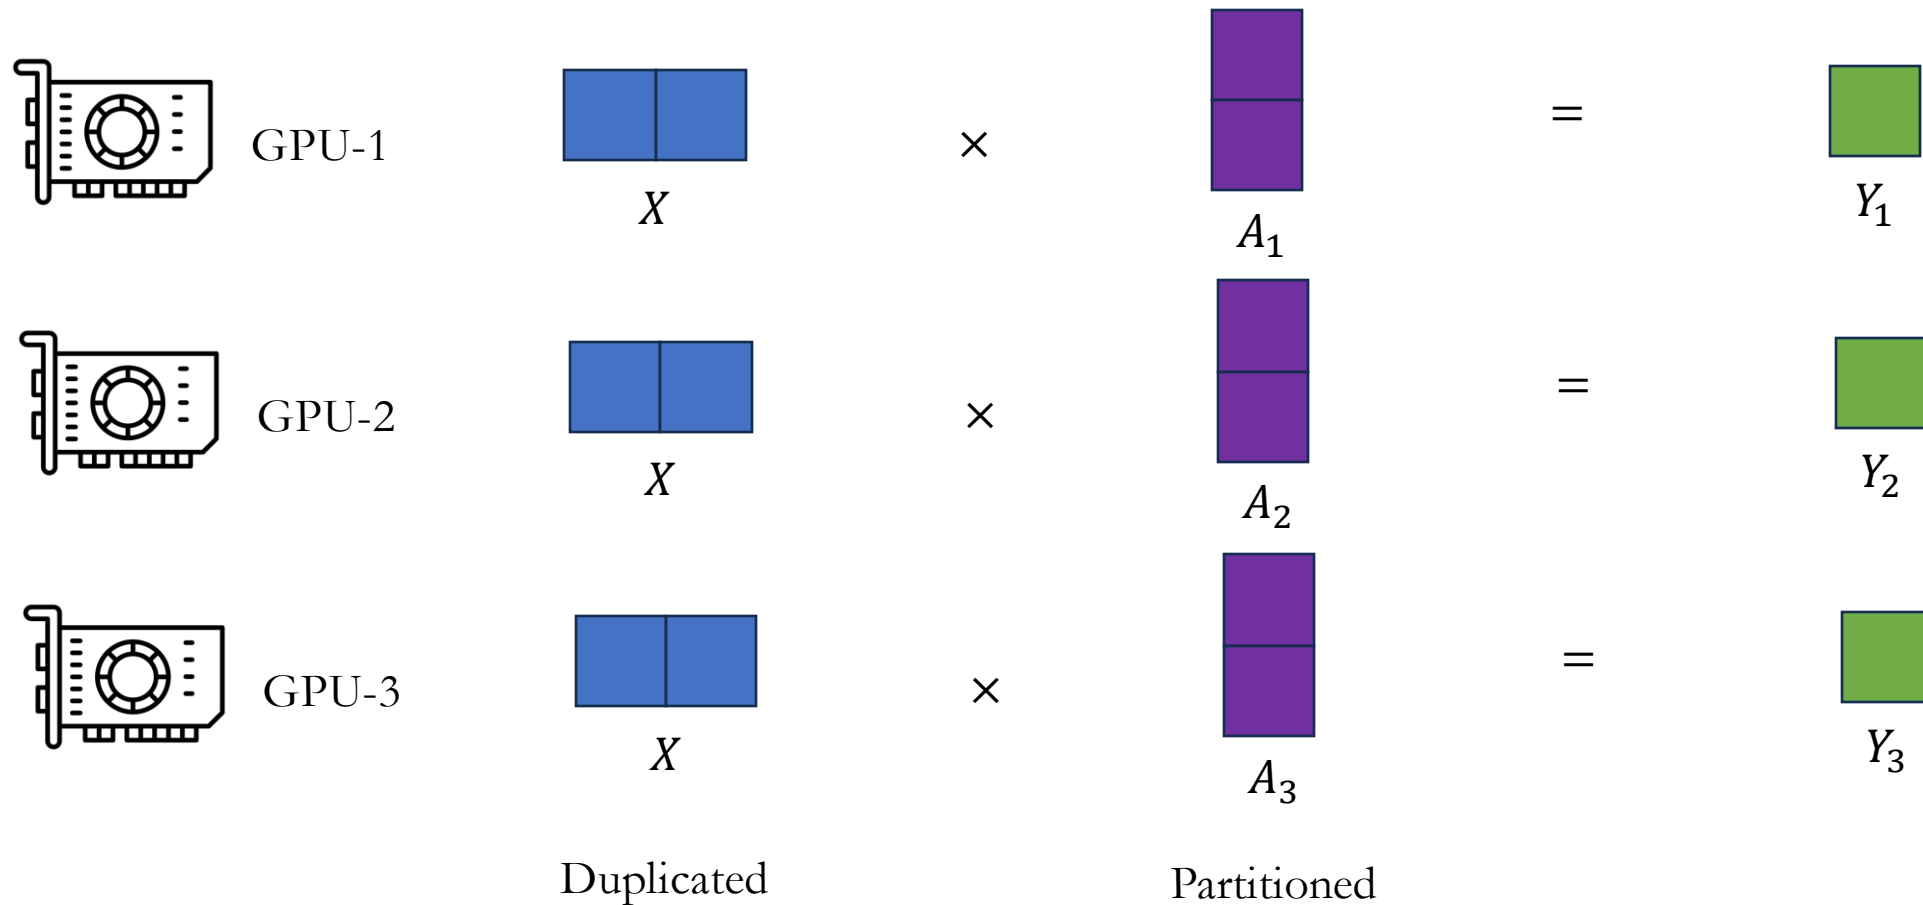
Partition the Matrix Multiplication

- Distribute the matrix multiplication:



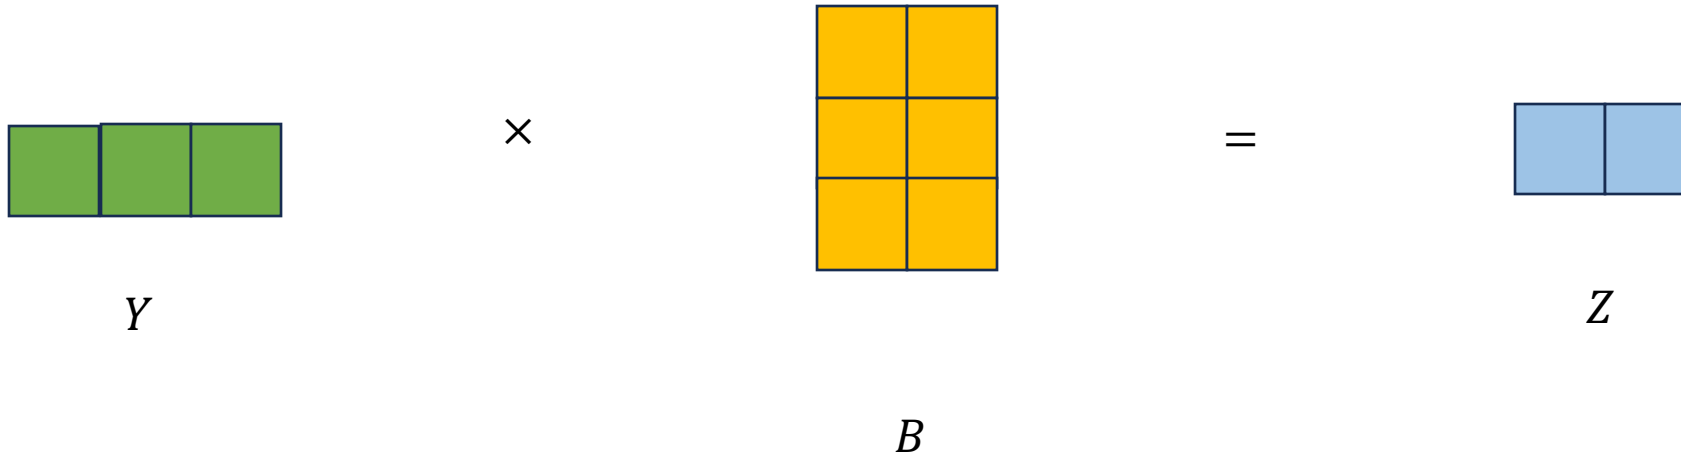
Partition the Matrix Multiplication

- Distribute the matrix multiplication:



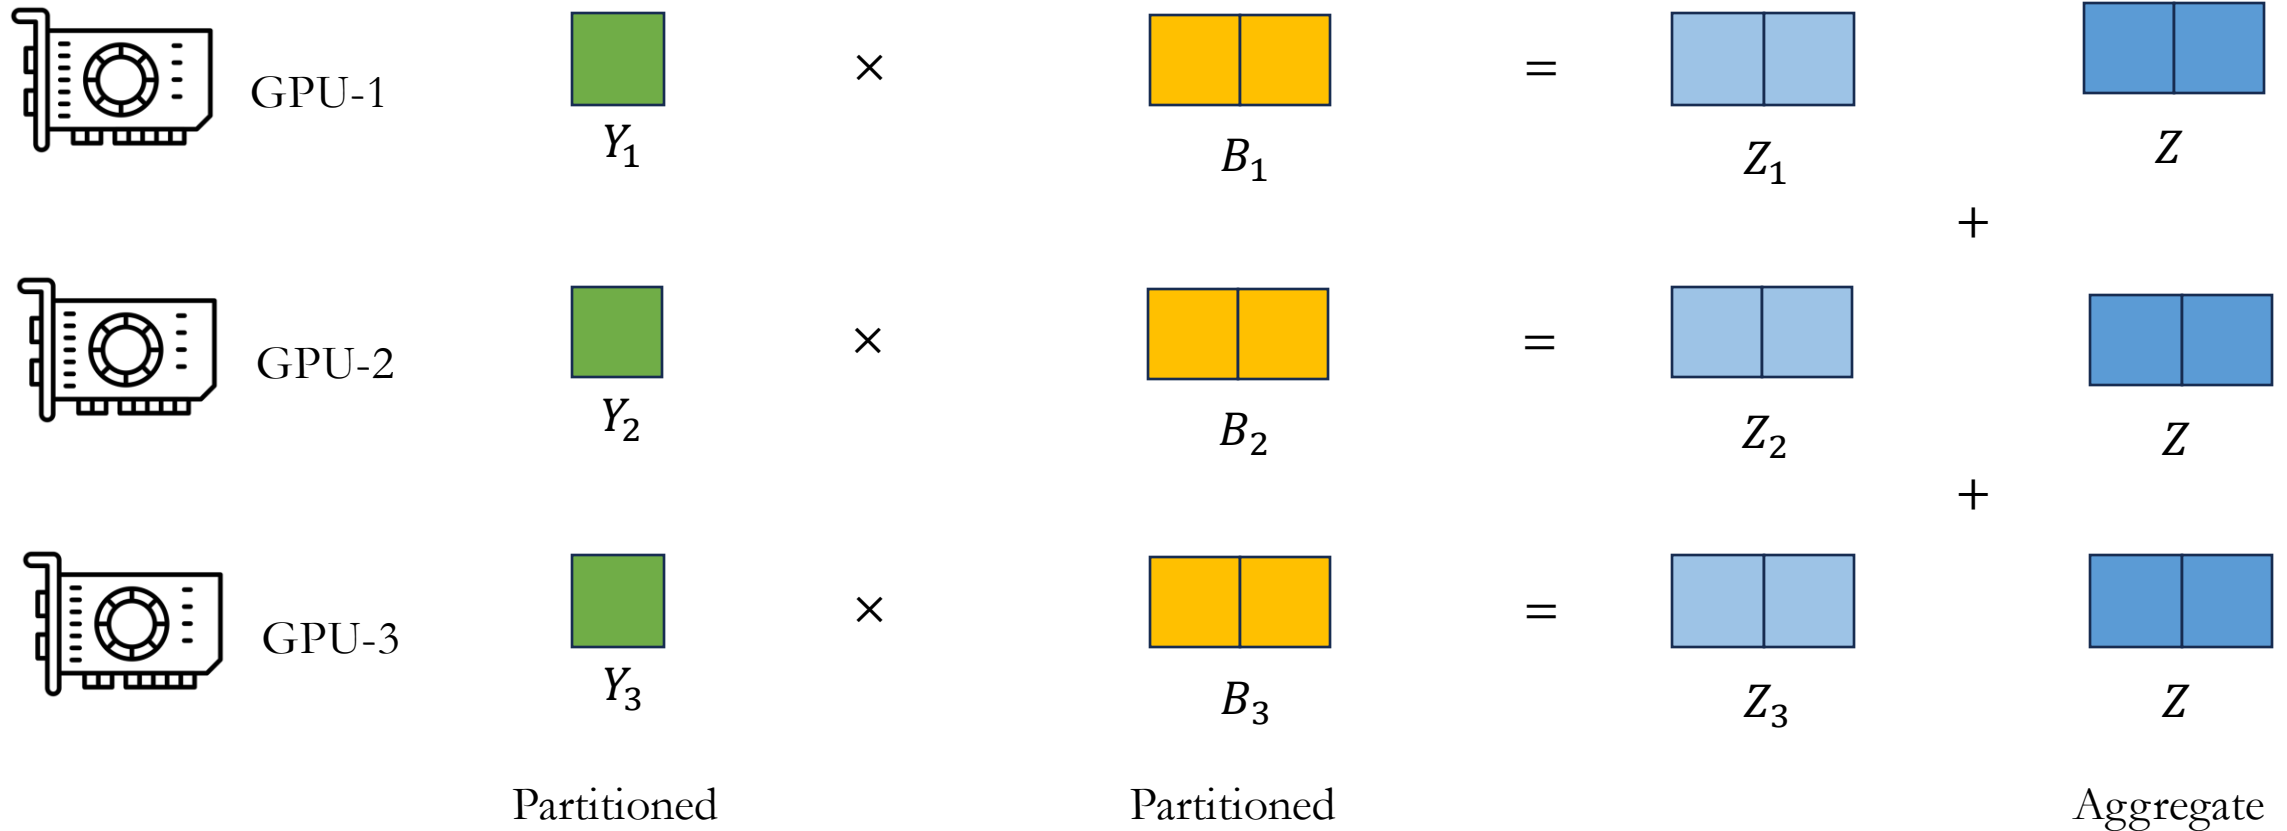
Partition the Matrix Multiplication

- Distribute the matrix multiplication:



Partition the Matrix Multiplication

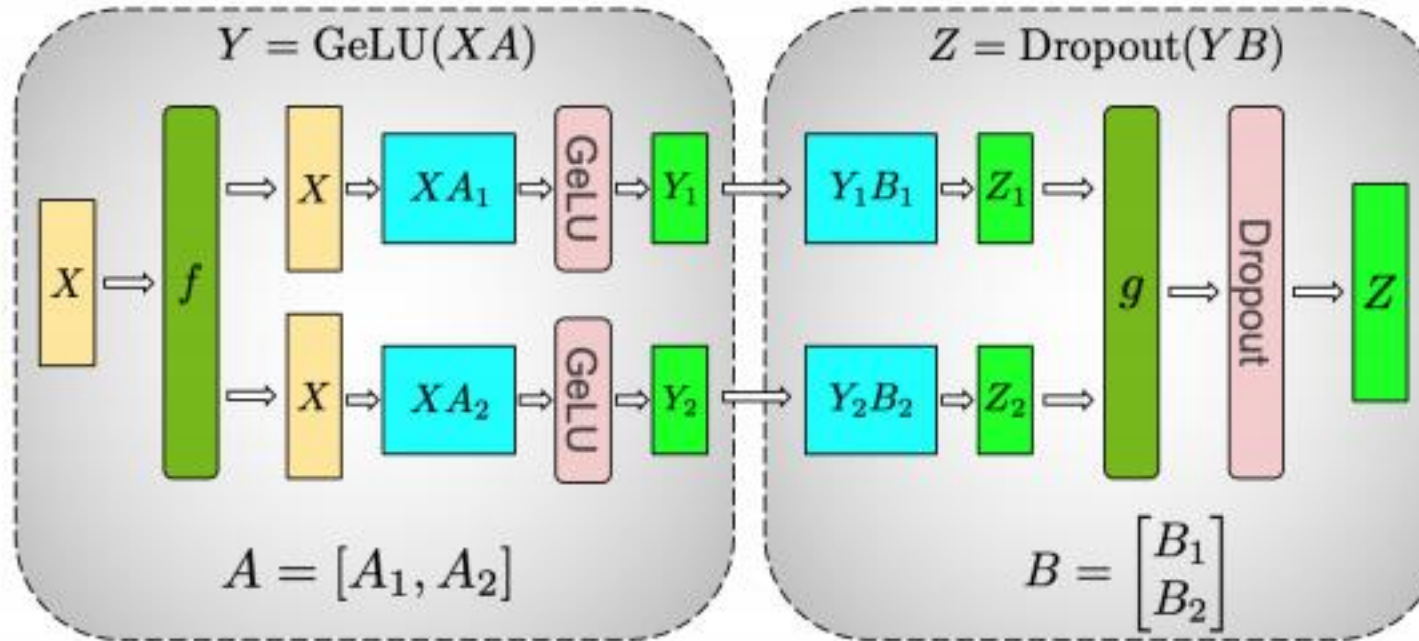
- Distribute the matrix multiplication:



Tensor Model Parallelism

- Split the first weight matrix col-wisely;
- Split the second weight matrix row-wisely;
- Duplicate the input on each GPU;
- Apply the computation as we illustrated above;
- Aggregate the outputs after the local computation.


MLP in Tensor Model Parallelism



(a) MLP

- f is the identity operator in the forward pass and the **AllReduce** operator in the backward pass.
- g is the **AllReduce** operator in the forward pass and the identity operator in the backward pass.

Multi-Head Attention in Tensor Model Parallelism



(b) Self-Attention

- f is the identity operator in the forward pass and the **AllReduce** operator in the backward pass.
- g is the **AllReduce** operator in the forward pass and the identity operator in the backward pass.

TransformerBlocks in Tensor Model Parallelism

- B is the batch size;
- L is the sequence length;
- D is the model dimension;
- Multi-head attention:
 $D = n_H \times H$
- H is the head dimension;
- n_H is the number of heads.
- d_{tp} is the tensor parallel degree: $d_{tp} \leq n_H$.

Computation	Input	Output
$Q = xW^Q$	$x \in \mathbb{R}^{B \times L \times D}, W^Q \in \mathbb{R}^{D \times \frac{D}{d_{tp}}}$	$Q \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$
$K = xW^K$	$x \in \mathbb{R}^{B \times L \times D}, W^K \in \mathbb{R}^{D \times \frac{D}{d_{tp}}}$	$K \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$
$V = xW^V$	$x \in \mathbb{R}^{B \times L \times D}, W^V \in \mathbb{R}^{D \times \frac{D}{d_{tp}}}$	$V \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$
$[Q_1, Q_2 \dots, Q_{\frac{n_H}{d_{tp}}}] = \text{Partion}_{-1}(Q)$	$Q \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$	$Q_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, \frac{n_H}{d_{tp}}$
$[K_1, K_2 \dots, K_{\frac{n_H}{d_{tp}}}] = \text{Partion}_{-1}(K)$	$K \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$	$K_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, \frac{n_H}{d_{tp}}$
$[V_1, V_2 \dots, V_{\frac{n_H}{d_{tp}}}] = \text{Partion}_{-1}(V)$	$V \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$	$V_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, \frac{n_H}{d_{tp}}$
$\text{Score}_i = \text{softmax}(\frac{Q_i K_i^T}{\sqrt{D}}), i = 1, \dots, \frac{n_H}{d_{tp}}$	$Q_i, K_i \in \mathbb{R}^{B \times L \times H}$	$\text{score}_i \in \mathbb{R}^{B \times L \times L}$
$Z_i = \text{score}_i V_i, i = 1, \dots, \frac{n_H}{d_{tp}}$	$\text{score}_i \in \mathbb{R}^{B \times L \times L}, V_i \in \mathbb{R}^{B \times L \times H}$	$Z_i \in \mathbb{R}^{B \times L \times H}$
$Z = \text{Merge}_{-1}([Z_1, Z_2 \dots, Z_{\frac{n_H}{d_{tp}}}])$	$Z_i \in \mathbb{R}^{B \times L \times H}, i = 1, \dots, \frac{n_H}{d_{tp}}$	$Z \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}$
$\text{Out} = ZW^O$	$Z \in \mathbb{R}^{B \times L \times \frac{D}{d_{tp}}}, W^O \in \mathbb{R}^{\frac{D}{d_{tp}} \times D}$	$\text{Out} \in \mathbb{R}^{B \times L \times D}$
AllReduce (Out)	$\text{Out} \in \mathbb{R}^{B \times L \times D}$	$\text{Out} \in \mathbb{R}^{B \times L \times D}$

TransformerBlocks in Tensor Model Parallelism

- B is the batch size;
- L is the sequence length;
- D is the model dimension;
- Multi-head attention:

$$D = n_H \times H$$
- H is the head dimension;
- n_H is the number of heads.
- d_{tp} is the tensor parallel degree: $d_{tp} \leq n_H$.

Computation	Input	Output
$A = \text{Out} W^1$	$\text{Out} \in \mathbb{R}^{B \times L \times D}, W^1 \in \mathbb{R}^{D \times \frac{4D}{d_{tp}}}$	$A \in \mathbb{R}^{B \times L \times \frac{4D}{d_{tp}}}$
$A' = \text{relu}(A)$	$A \in \mathbb{R}^{B \times L \times \frac{4D}{d_{tp}}}$	$A' \in \mathbb{R}^{B \times L \times \frac{4D}{d_{tp}}}$
$x' = A' W^2$	$A' \in \mathbb{R}^{B \times L \times \frac{4D}{d_{tp}}}, W^2 \in \mathbb{R}^{\frac{4D}{d_{tp}} \times D}$	$x' \in \mathbb{R}^{B \times L \times D}$
AllReduce (x')	$x' \in \mathbb{R}^{B \times L \times D}$	$x' \in \mathbb{R}^{B \times L \times D}$



Tensor Model Parallelism in Megatron-LM

<https://github.com/NVIDIA/Megatron-LM>

Entrance of the Training Scripts

```
model = GPTModel(  
    config=config,  
    transformer_layer_spec=transformer_layer_spec,  
    vocab_size=args.padded_vocab_size,  
    max_sequence_length=args.max_position_embeddings,  
    pre_process=pre_process,  
    post_process=post_process,  
    fp16_lm_cross_entropy=args.fp16_lm_cross_entropy,  
    parallel_output=True,  
    share_embeddings_and_output_weights=not args.untie_embeddings_and_output_weights,  
    position_embedding_type=args.position_embedding_type,  
    rotary_percent=args.rotary_percent,  
)
```

https://github.com/NVIDIA/Megatron-LM/blob/main/pretrain_gpt.py#L60

Launch Scripts

```
GPT_ARGS="  
  --tensor-model-parallel-size 2 \  
  --pipeline-model-parallel-size 2 \  
  --sequence-parallel \  
  --num-layers 24 \  
  --hidden-size 1024 \  
  --num-attention-heads 16 \  
  --seq-length 1024 \  
  --max-position-embeddings 1024 \  
  --micro-batch-size 4 \  
  --global-batch-size 16 \  
  --lr 0.00015 \  
  --train-iters 500000 \  
  --lr-decay-iters 320000 \  
  --lr-decay-style cosine \  
  --min-lr 1.0e-5 \  
  --weight-decay 1e-2 \  
  --lr-warmup-fraction .01 \  
  --clip-grad 1.0 \  
  --fp16  
"
```

References

- <https://huggingface.co/transformers/v4.9.2/parallelism.html>
- <https://arxiv.org/abs/2104.04473>
- <https://github.com/NVIDIA/Megatron-LM/tree/main>