



Parameter Efficient Fine-Tuning

COMP4901Y

Binhang Yuan

Fine-Tuning

How LLM Are Usually Deployed?


- **Pre-training** is the initial phase of training an LLM, during which it learns from a large, diverse dataset, often consisting of trillions of tokens.
 - The goal is to develop a broad understanding of language, context, and various types of knowledge for the model.
 - Pre-training is usually computationally intensive (thousands of GPUs for weeks) and requires huge amounts of data (trillions of tokens).
- **Fine-tuning** is where you take an already pre-trained model and further train it on a more specific dataset.
 - This dataset is typically smaller and focused on a particular domain or task.
 - The purpose of fine-tuning is to adapt the model to perform better in specific scenarios or on tasks that were not well covered during pre-training.
 - The new knowledge added during fine-tuning enhances the model's performance in specific contexts rather than broadly expanding its general knowledge.

How LLM Are Usually Deployed?



Stage 1: Pretraining

1. Prepare 10TB of text as the training corpus.
2. Use a cluster of thousands of GPUs to train a neural network with the corpus from scratch.
3. Obtain the *base model*.



Stage 2: Fine-tuning

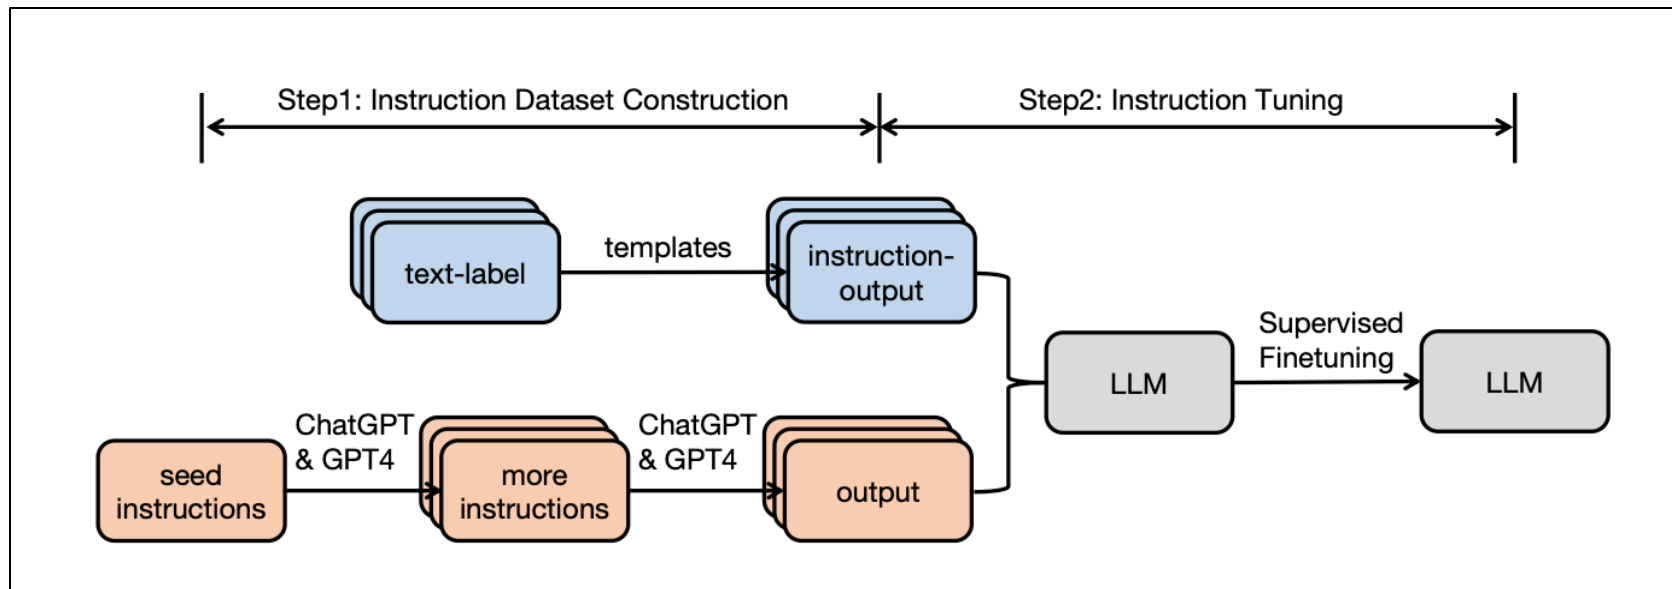
1. Prepare 100K high-quality text as the task-specific training dataset.
2. Use a small-scale GPU cluster to fine-tune the base model with the training dataset.
3. Obtain the *domain-specific model*.
4. Run the domain-specific model for evaluation.
5. Repeat until we are happy with the results.

Instruction Tuning

- Fine-tune the LLM using (Instruction, output) pairs:
 - Instruction: denotes the human instruction for the model;
 - Output: denotes the desired output that follows the instruction.
- Finetuning an LLM on the instruction dataset bridges the gap between the next-word prediction objective of LLMs and the users' objective of instruction following;
- Instruction tuning allows for a more controllable and predictable model behavior compared to standard LLMs.
 - The instructions serve to constrain the model's outputs to align with the desired response characteristics or domain knowledge, providing a channel for humans to intervene with the model's behaviors.
- Instruction tuning can help LLMs rapidly adapt to a specific domain without extensive retraining or architectural changes.

Instruction Tuning

- Two methods to construct the instruction datasets:
 - Data integration from annotated natural language datasets: (instruction, output) pairs are collected from existing annotated natural language datasets by using templates to transform text-label pairs to (instruction, output) pairs.
 - Generate outputs using more advanced LLMs: employ LLMs such as GPT-3.5-Turbo or GPT4 to gather the desired outputs given the instructions instead of manually collecting the outputs.



Fine-Tuning v.s. Prompt Engineering

- Suppose we have:
 - A dataset $D = \{(x_i, y_i)\}_{i=1}^N$ and N is rather small.
 - A pre-trained LLM.
- How to fit it to your task?

- Option A: Fine-tuning:
 - Fine-tune the LLM on the training data using:
 - A standard training objective;
 - SGD to update (part of) the LLM's parameters.
 - Advantages:
 - Fits into the standard ML recipe;
 - Still works if N becomes relatively large.
 - Disadvantages:
 - Backward pass is computationally expensive in terms of FLOPs and memory footprint;
 - You have to have full access of the pre-trained LLM.

- Option B: Prompt engineering (in-context learning):
 - Feed training examples to the LLM as a prompt:
 - Allow the LLM to infer patterns in the training examples during inference;
 - Take the output of the LLM following the prompt as its prediction.
 - Advantages:
 - No backpropagation required and only one pass through the training data;
 - Does not require model weights, only API access.
 - Disadvantages:
 - The prompt may be very long.

Fine-Tuning v.s. Prompt Engineering

- Why would we ever bother with fine-tuning if it's so inefficient?
 - Because, even for very large LMs, fine-tuning often beats in-context learning.
 - In a fair comparison of fine-tuning (FT) and in-context learning (ICL), we find that FT outperforms ICL for most model sizes.

		FT									FT						
		125M	350M	1.3B	2.7B	6.7B	13B	30B			125M	350M	1.3B	2.7B	6.7B	13B	30B
ICL	125M	-0.00	0.01	0.02	0.03	0.12	0.14	0.09	ICL	125M	-0.00	0.00	0.02	0.01	0.10	0.11	0.07
	350M	-0.00	0.01	0.02	0.03	0.12	0.14	0.09		350M	-0.00	0.00	0.02	0.01	0.10	0.11	0.07
	1.3B	-0.00	0.01	0.02	0.03	0.12	0.14	0.09		1.3B	-0.01	-0.00	0.01	0.01	0.10	0.11	0.07
	2.7B	-0.00	0.01	0.02	0.03	0.12	0.14	0.09		2.7B	-0.01	-0.00	0.01	0.01	0.09	0.10	0.07
	6.7B	-0.00	0.01	0.02	0.03	0.12	0.14	0.09		6.7B	-0.01	-0.01	0.01	0.00	0.09	0.10	0.06
	13B	-0.04	-0.02	-0.01	-0.00	0.09	0.11	0.05		13B	-0.03	-0.03	-0.02	-0.02	0.07	0.08	0.04
	30B	-0.11	-0.09	-0.08	-0.08	0.02	0.03	-0.02		30B	-0.07	-0.07	-0.05	-0.06	0.03	0.04	0.00

(a) RTE

(b) MNLI

Table 1: Difference between average **out-of-domain performance** of ICL and FT on RTE (a) and MNLI (b) across model sizes. We use 16 examples and 10 random seeds for both approaches. For ICL, we use the `gpt-3` pattern. For FT, we use pattern-based fine-tuning (PBFT) and select checkpoints according to in-domain performance. We perform a Welch's t-test and color cells according to whether: **ICL performs significantly better than FT**, **FT performs significantly better than ICL**. For cells without color, there is no significant difference.

Parameter Efficient Fine-Tuning

Parameter Efficient Fine-Tuning


- Parameter efficient fine-tuning (PEFT): Rather than finetuning the entire model, we finetune only small amounts of weights.
- Goal: achieve performance on a downstream task that is comparable to fine-tuning all parameters.
- Some approaches:
 - Frozen layer/Subset fine-tuning: pick a subset of the parameters, fine-tune only those layers, and freeze the rest of the layers.
 - Adapters: add additional layers that have few parameters and tune only the parameters of those layers, keeping all others fixed.
 - Low-rank adaption (LoRA): learn a low rank approximation of the weight matrices.

Subset Fine-Tuning

- Some interpretations from NLP research:
 - Earlier layers of the transformer tend to capture linguistic phenomena and basic language understanding;
 - Later layers are where the task-specific learning happens.
- We should be able to learn new tasks by freezing the earlier layers and tuning the later ones.
- This can be a simple baseline for PEFT.

Subset Fine-Tuning

- Keep all parameters fixed except for the top K layers.
- Gradients only need to flow through top K layers instead of all of layers.
- Reduce computation load.
- Reduce memory usage.

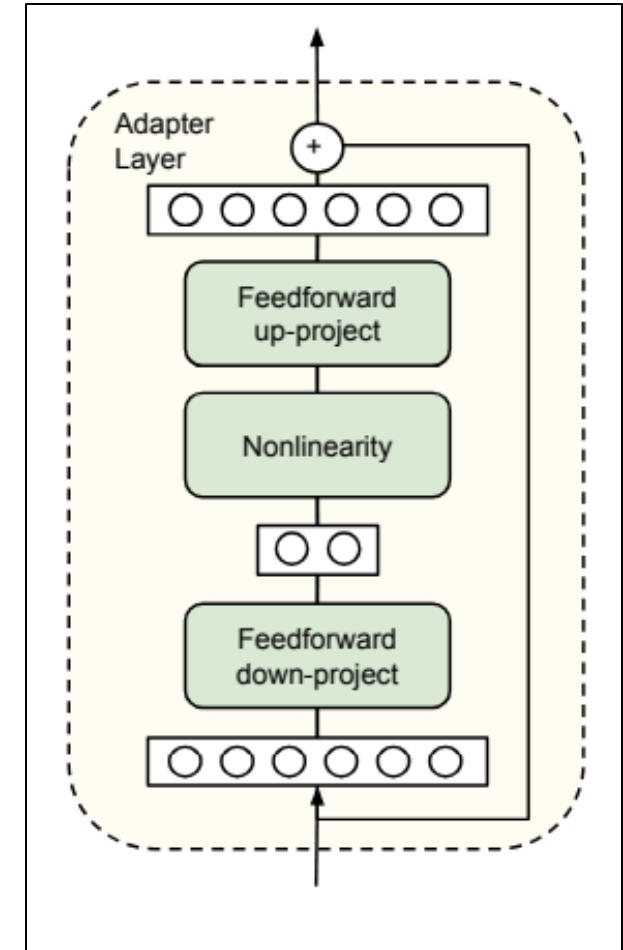


The gradient does not backpropagate to lower layers.

Adapter

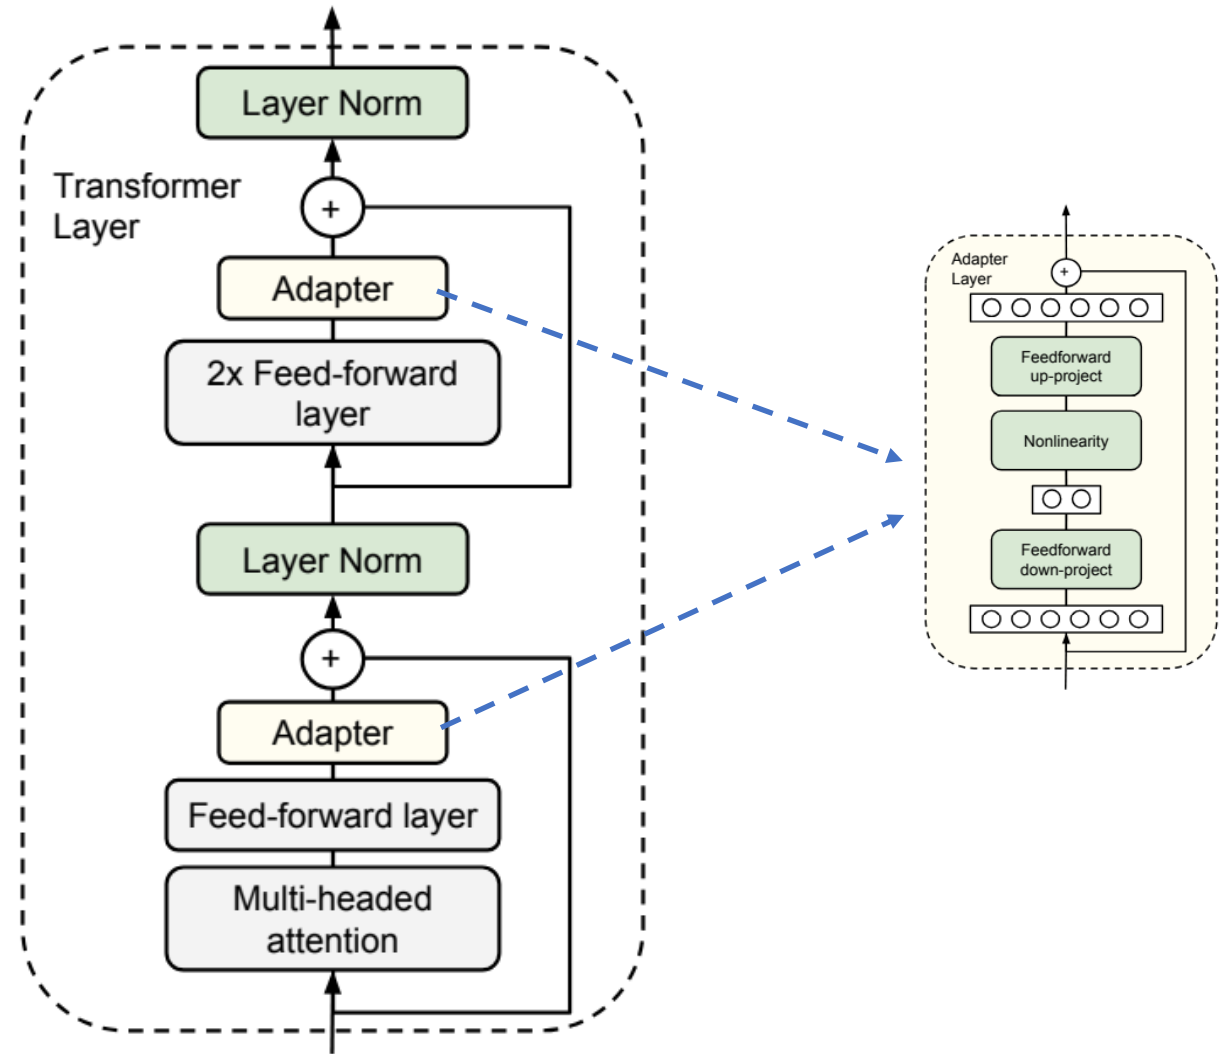


- Adapters are new modules are added between layers of a pre-trained network.
 - The original model weights are fixed;
 - Just the adapter modules are tuned.
- An adapter layer is simply a feed- forward neural network with one hidden layer, and a residual connection.
- Suppose the original LLM has a model dimension of D :
 - For the down-project weight matrix: $W_A \in \mathbb{R}^{D \times R}$;
 - For the up-project weight matrix: $W_B \in \mathbb{R}^{R \times D}$;
 - We have $R \ll D$.



Adapter

- Given $R \ll D$, in practice the adapter layers contain only 0.5% – 8% of the total parameters.
- When added to a deep neural network (e.g. transformer) all the other parameters of the pretrained model are kept fixed, and only the adapter layer parameters are fine-tuned.



Low-Rank Adaptation (LoRA)

Low-Rank Adaption

- Central idea:
 - “How can we re-parameterize the model into something more efficient to train?”
- Finetuning has a low intrinsic dimension, that is, the minimum number of parameters needed to be modified to reach satisfactory performance is not very large.
- This means we can re-parameterize a subset of the original model parameters with low-dimensional proxy parameters, and just optimize the proxy.

Intrinsic Dimension

- An objective function's intrinsic dimension measures the minimum number of parameters needed to reach a satisfactory solution to the objective.
- Can also be thought of as the lowest dimensional subspace in which one can optimize the original objective function to within a certain level of approximation error.
- Details in this paper: <https://arxiv.org/abs/2012.13255>
 - Suppose we have model parameters $\theta^{(D)} \in \mathbb{R}^D$, D is the number of parameters;
 - Instead of optimizing $\theta^{(D)}$, we could instead optimize a smaller set of parameters $\theta^{(d)} \in \mathbb{R}^d$, where $d \ll D$.
 - This done through clever factorization:
 - $\theta^{(D)} = \theta_0^{(D)} + P(\theta^{(d)})$; where $P: \mathbb{R}^d \rightarrow \mathbb{R}^D$
 - P is typically a linear projection: $\theta^{(D)} = \theta_0^{(D)} + \theta^{(d)} M$.
 - Intuitively, we do an arbitrary random projection onto a much smaller space; usually, a linear projection, we then solve the optimization problem in that smaller subspace. If we reach a satisfactory solution, we say the dimensionality of that subspace is the intrinsic dimension.

Low-Rank Adaption


- Full paper: <https://arxiv.org/pdf/2106.09685>.
- Intuition: It's not just the model weights that are low rank, updates to the model weights are also low-rank.
- LoRA freezes the pre-trained model weights and injects trainable rank decomposition matrices into some or all layers.

LoRA Key Idea

- Keep the original pre-trained parameters W fixed during fine-tuning;
- Learn an additive modification to those parameters ΔW ;
- Define ΔW as a low-rank decomposition: $\Delta W = AB$.


Standard Linear Layer

- $Y = XW$
- $X \in \mathbb{R}^{D_1}, Y \in \mathbb{R}^{D_2}, W \in \mathbb{R}^{D_1 \times D_2}$



LoRA Linear Layer

- $Y = XW + XAB = X(W + AB)$
- $X \in \mathbb{R}^{D_1}, Y \in \mathbb{R}^{D_2}, W \in \mathbb{R}^{D_1 \times D_2}$
- $A \in \mathbb{R}^{D_1 \times R}, B \in \mathbb{R}^{R \times D_2}$




LoRA Initialization

- We initialize the trainable parameters:
 - $A_{ij} \sim \mathcal{N}(0, \sigma^2)$ $B_{ij} = 0$
- This ensures that, at the start of fine-tuning, the parameters have their pre-trained values:
 - $\Delta W = AB = 0$

LoRA Linear Layer

- $Y = XW + XAB = X(W + AB)$
- $X \in \mathbb{R}^{D_1}, Y \in \mathbb{R}^{D_2}, W \in \mathbb{R}^{D_1 \times D_2}$
- $A \in \mathbb{R}^{D_1 \times R}, B \in \mathbb{R}^{R \times D_2}$




LoRA Hot Swapping Parameters

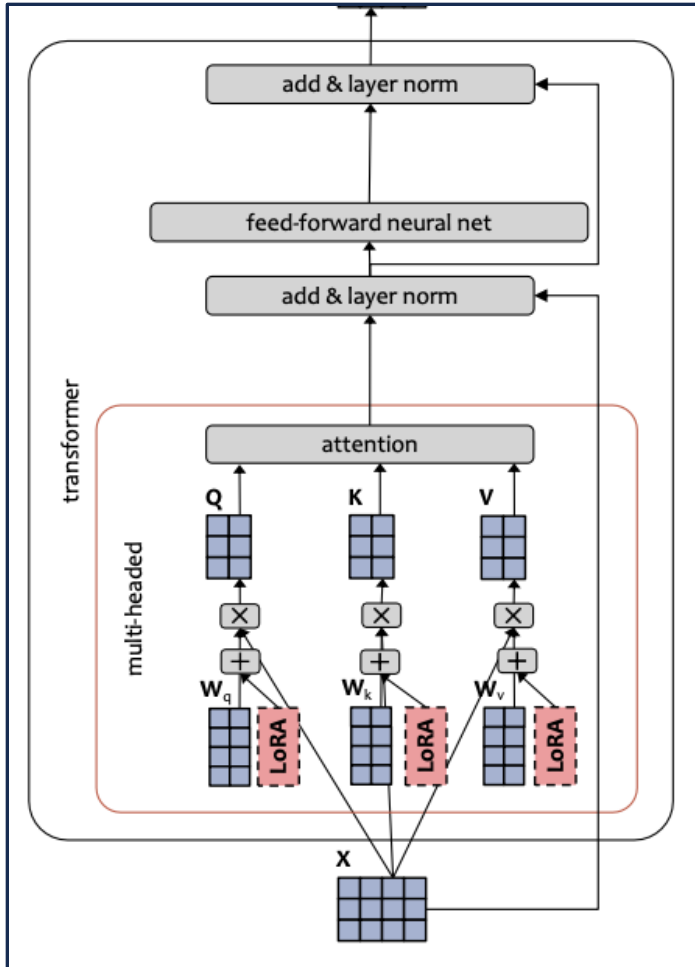
- W and AB the same dimension, so we can swap the LoRA parameters in and out of a Standard Linear Layer.
- To include LoRA:
 - $W' \leftarrow W + AB$
- To remove LoRA:
 - $W \leftarrow W' - AB$

LoRA Linear Layer

- $Y = XW + XAB = X(W + AB)$
- $X \in \mathbb{R}^{D_1}, Y \in \mathbb{R}^{D_2}, W \in \mathbb{R}^{D_1 \times D_2}$
- $A \in \mathbb{R}^{D_1 \times R}, B \in \mathbb{R}^{R \times D_2}$



LoRA for Transformers



- LoRA linear layers could replace every linear layer in the Transformer layer;
- But the original paper only applies LoRA to the attention weights:
 - $Q = \text{LoRALinear}(X, W_q, A_q, B_q)$
 - $K = \text{LoRALinear}(X, W_k, A_k, B_k)$
 - $V = \text{LoRALinear}(X, W_v, A_v, B_v)$
- Some further research found that most efficient to include LoRA only on the query and key linear layers.

LoRA Results

- Some empirical takeaways:
 - Applied to GPT-3, LoRA achieves performance almost as good as full parameter fine-tuning, but with far fewer parameters.
 - On some tasks, it even outperforms full fine-tuning.
 - For some datasets, a rank of $R = 1$ is sufficient.
 - LoRA performs well when the dataset is large or small.

References

- https://www.youtube.com/watch?v=zjkBMFhNj_g
- <https://arxiv.org/abs/2308.10792>
- https://www.andrew.cmu.edu/course/11-667/lectures/W4L2_PETM.pptx.pdf
- <https://www.cs.cmu.edu/~mgormley/courses/10423//slides/lecture11-peft-ink.pdf>
- <https://arxiv.org/abs/2012.13255>
- <https://arxiv.org/pdf/2106.09685>